

Lead Scoring Case Study

Problem Statement

An education company named X Education sells online courses to industry professionals. On any given day, many professionals who are interested in the courses land on their website and browse for courses.

The company markets its courses on several websites and search engines like Google. Once these people land on the website, they might browse the courses or fill up a form for the course or watch some videos. When these people fill up a form providing their email address or phone number, they are classified to be a lead. Moreover, the company also gets leads through past referrals. Once these leads are acquired, employees from the sales team start making calls, writing emails, etc. Through this process, some of the leads get converted while most do not. The typical lead conversion rate at X education is around 30%.

Now, although X Education gets a lot of leads, its lead conversion rate is very poor. For example, if, say, they acquire 100 leads in a day, only about 30 of them are converted. To make this process more efficient, the company wishes to identify the most potential leads, also known as 'Hot Leads'. If they successfully identify this set of leads, the lead conversion rate should go up as the sales team will now be focusing more on communicating with the potential leads rather than making calls to everyone.

There are a lot of leads generated in the initial stage (top) but only a few of them come out as paying customers from the bottom. In the middle stage, we need to nurture the potential leads well (i.e. educating the leads about the product, constantly communicating etc.) in order to get a higher lead conversion.

X Education has appointed me to help them select the most promising leads, i.e. the leads that are most likely to convert into paying customers. The company requires me to build a model wherein I'll need to assign a lead score to each of the leads such that the customers with higher lead score have a higher conversion chance and the customers with lower lead score have a lower conversion chance. The CEO, in particular, has given a ballpark of the target lead conversion rate to be around 80%.

Data

You have been provided with a leads dataset from the past with around 9000 data points. This dataset consists of various attributes such as Lead Source, Total Time Spent on Website, Total Visits, Last Activity, etc. which may or may not be useful in ultimately deciding whether a lead will be converted or not. The target variable, in this case, is the column 'Converted' which tells whether a past lead was converted or not wherein 1 means it was converted and 0 means it wasn't converted. You can learn more about the dataset from the data dictionary provided in the zip folder at the end of the page. Another thing that you also need to check out for are the levels present in the categorical variables. Many of the categorical variables have a level called 'Select' which needs to be handled because it is as good as a null value (think why?).

Table of Contents

- [1 Reading and Understanding Data](#)
- [2 Data Cleaning](#)
 - [2.1 Rename column names](#)
 - [2.2 Drop prospect_id column](#)
 - [2.3 Replace "Select" category with null values](#)

- [2.4 Handle null values and sales generated columns](#)
 - [2.4.1 Drop columns that have null values > 40% or Sales generated columns](#)
 - [2.4.2 country column](#)
 - [2.4.3 course_selection_reason column](#)
 - [2.4.4 occupation column](#)
 - [2.4.5 specialization column](#)
 - [2.4.6 city column](#)
 - [2.5 Handle categorical columns with low number of missing values and low representation of categories](#)
 - [2.5.1 lead_origin column](#)
 - [2.5.2 lead_source column](#)
 - [2.6 Handle Binary columns](#)
 - [2.7 Handle Numerical columns](#)
 - [2.7.1 lead_number column: change datatype](#)
 - [2.7.2 total_visits column](#)
 - [2.7.3 page_views_per_visit column](#)
 - [3 Exploratory Data Analysis](#)
 - [3.1 Numerical columns](#)
 - [3.1.1 Heatmap](#)
 - [3.1.2 Check for outliers](#)
 - [3.2 Categorical columns](#)
 - [3.2.1 Lead Origin](#)
 - [3.2.2 Lead Source](#)
 - [3.2.3 Specialization](#)
 - [3.2.4 Occupation](#)
 - [3.2.5 City](#)
 - [4 Data Preparation](#)
 - [4.1 Converting Binary \(Yes/No\) to 0/1](#)
 - [4.2 Creating dummy variable for categorical columns](#)
 - [4.3 Outliers Treatment](#)
 - [4.4 Test-Train Split](#)
 - [4.5 Feature Scaling](#)
 - [4.6 Looking at correlations](#)
 - [4.6.1 Drop highly correlated dummy variables](#)
 - [5 Model Building](#)
 - [5.1 Model 1: All variables](#)
 - [5.2 Feature selection using RFE](#)
 - [5.3 Model 2: Assessing the model with statsmodel](#)
-

```

# basic libraries to work on the dataframe
import pandas as pd
import numpy as np
# data Visualization libraries
import matplotlib.pyplot as plt
import seaborn as sns

# libraries
from sklearn.cluster import KMeans
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import statsmodels.api as sm
from sklearn.linear_model import LogisticRegression
from sklearn.feature_selection import RFE
from statsmodels.stats.outliers_influence import variance_inflation_factor

# Suppressing Warnings
import warnings
warnings.filterwarnings('ignore')

#Increasing the columns views limit
pd.options.display.max_columns = None
pd.options.display.max_rows = 150
pd.options.display.float_format = '{:.2f}'.format

```

Reading and Understanding Data

```

#Reading the data file using pandas
df = pd.read_csv('Leads.csv')

df.head()

```

| | Prospect ID | Lead Number | Lead Origin | Lead Source | Do Not Email | Do Not Call | Converted | TotalVisits | Total Time Spent on Website | Page Views Per Visit | Last Activity |
|---|--------------------------------------|-------------|-------------------------|----------------|--------------|-------------|-----------|-------------|-----------------------------|----------------------|-------------------------|
| 0 | 7927b2df-8bba-4d29-b9a2-b6e0beafe620 | 660737 | API | Olark Chat | No | No | 0 | 0.00 | 0 | 0.00 | Page Visited on Website |
| 1 | 2a272436-5132-4136-86fa-dcc88c88f482 | 660728 | API | Organic Search | No | No | 0 | 5.00 | 674 | 2.50 | Email Opened |
| 2 | 8cc8c611-a219-4f35-ad23-fdfd2656bd8a | 660727 | Landing Page Submission | Direct Traffic | No | No | 1 | 2.00 | 1532 | 2.00 | Email Opened |
| 3 | 0cc2df48-7cf4-4e39-9de9-19797f9b38cc | 660719 | Landing Page Submission | Direct Traffic | No | No | 0 | 1.00 | 305 | 1.00 | Unreachable |

| | Prospect ID | Lead Number | Lead Origin | Lead Source | Do Not Email | Do Not Call | Converted | TotalVisits | Total Time Spent on Website | Page Views Per Visit | Last Activity |
|---|--------------------------------------|-------------|-------------------------|-------------|--------------|-------------|-----------|-------------|-----------------------------|----------------------|-------------------|
| 4 | 3256f628-e534-4826-9d63-4a8b88782852 | 660681 | Landing Page Submission | Google | No | No | 1 | 2.00 | 1428 | 1.00 | Converted to Lead |

```
# check the shape of the dataset
df.shape
```

```
(9240, 37)
```

```
# check statistics for numerical columns
df.describe()
```

| | Lead Number | Converted | TotalVisits | Total Time Spent on Website | Page Views Per Visit | Asymmetrique Activity Score | Asymmetrique Profile Score |
|-------|-------------|-----------|-------------|-----------------------------|----------------------|-----------------------------|----------------------------|
| count | 9240.00 | 9240.00 | 9103.00 | 9240.00 | 9103.00 | 5022.00 | 5022.00 |
| mean | 617188.44 | 0.39 | 3.45 | 487.70 | 2.36 | 14.31 | 16.34 |
| std | 23406.00 | 0.49 | 4.85 | 548.02 | 2.16 | 1.39 | 1.81 |
| min | 579533.00 | 0.00 | 0.00 | 0.00 | 0.00 | 7.00 | 11.00 |
| 25% | 596484.50 | 0.00 | 1.00 | 12.00 | 1.00 | 14.00 | 15.00 |
| 50% | 615479.00 | 0.00 | 3.00 | 248.00 | 2.00 | 14.00 | 16.00 |
| 75% | 637387.25 | 1.00 | 5.00 | 936.00 | 3.00 | 15.00 | 18.00 |
| max | 660737.00 | 1.00 | 251.00 | 2272.00 | 55.00 | 18.00 | 20.00 |

```
# check whether there are any duplicates
df.duplicated().sum()
```

```
0
```

```
#Lets have a look at all the columns, their datatypes and also get an idea of null values
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 9240 entries, 0 to 9239
```

```
Data columns (total 37 columns):
```

| # | Column | Non-Null Count | Dtype |
|-----|--------------|----------------|--------|
| --- | ----- | ----- | ----- |
| 0 | Prospect ID | 9240 non-null | object |
| 1 | Lead Number | 9240 non-null | int64 |
| 2 | Lead Origin | 9240 non-null | object |
| 3 | Lead Source | 9204 non-null | object |
| 4 | Do Not Email | 9240 non-null | object |

| | | | | |
|----|---|------|----------|---------|
| 5 | Do Not Call | 9240 | non-null | object |
| 6 | Converted | 9240 | non-null | int64 |
| 7 | TotalVisits | 9103 | non-null | float64 |
| 8 | Total Time Spent on Website | 9240 | non-null | int64 |
| 9 | Page Views Per Visit | 9103 | non-null | float64 |
| 10 | Last Activity | 9137 | non-null | object |
| 11 | Country | 6779 | non-null | object |
| 12 | Specialization | 7802 | non-null | object |
| 13 | How did you hear about X Education | 7033 | non-null | object |
| 14 | What is your current occupation | 6550 | non-null | object |
| 15 | What matters most to you in choosing a course | 6531 | non-null | object |
| 16 | Search | 9240 | non-null | object |
| 17 | Magazine | 9240 | non-null | object |
| 18 | Newspaper Article | 9240 | non-null | object |
| 19 | X Education Forums | 9240 | non-null | object |
| 20 | Newspaper | 9240 | non-null | object |
| 21 | Digital Advertisement | 9240 | non-null | object |
| 22 | Through Recommendations | 9240 | non-null | object |
| 23 | Receive More Updates About Our Courses | 9240 | non-null | object |
| 24 | Tags | 5887 | non-null | object |
| 25 | Lead Quality | 4473 | non-null | object |
| 26 | Update me on Supply Chain Content | 9240 | non-null | object |
| 27 | Get updates on DM Content | 9240 | non-null | object |
| 28 | Lead Profile | 6531 | non-null | object |
| 29 | City | 7820 | non-null | object |
| 30 | Asymmetrique Activity Index | 5022 | non-null | object |
| 31 | Asymmetrique Profile Index | 5022 | non-null | object |
| 32 | Asymmetrique Activity Score | 5022 | non-null | float64 |
| 33 | Asymmetrique Profile Score | 5022 | non-null | float64 |
| 34 | I agree to pay the amount through cheque | 9240 | non-null | object |
| 35 | A free copy of Mastering The Interview | 9240 | non-null | object |
| 36 | Last Notable Activity | 9240 | non-null | object |

dtypes: float64(4), int64(3), object(30)

memory usage: 2.6+ MB

Observations

- A large number of columns have null values. Those columns should ideally be dropped
- Prospect ID and Lead Number both serve the same purpose. They are both unique identifiers. We will drop Prospect ID
- Column names are just too long. We will modify the column names
- Few categorical columns have "Select" in their entries. Those select are essentially null values because Select appears when someone does not select anything from the dropdown

Data Cleaning

Rename column names

- Long column names make analysis tiring as one has to always refer to column names. Also has impact on charts created later on
- Ideally, we should follow python's preferred Snakecase nomenclature

```
# change nomenclature to snakecase
df.columns = df.columns.str.replace(' ', '_').str.lower()

# test
df.columns
```

```
Index(['prospect_id', 'lead_number', 'lead_origin', 'lead_source',
      'do_not_email', 'do_not_call', 'converted', 'totalvisits',
      'total_time_spent_on_website', 'page_views_per_visit', 'last_activity',
      'country', 'specialization', 'how_did_you_hear_about_x_education',
      'what_is_your_current_occupation',
      'what_matters_most_to_you_in_choosing_a_course', 'search', 'magazine',
      'newspaper_article', 'x_education_forums', 'newspaper',
      'digital_advertisement', 'through_recommendations',
      'receive_more_updates_about_our_courses', 'tags', 'lead_quality',
      'update_me_on_supply_chain_content', 'get_updates_on_dm_content',
      'lead_profile', 'city', 'asymmetrique_activity_index',
      'asymmetrique_profile_index', 'asymmetrique_activity_score',
      'asymmetrique_profile_score',
      'i_agree_to_pay_the_amount_through_cheque',
      'a_free_copy_of_mastering_the_interview', 'last_notable_activity'],
      dtype='object')
```

```
# shorten column names
df.rename(columns = {'totalvisits': 'total_visits', 'total_time_spent_on_website': 'time_on_website',
                    'how_did_you_hear_about_x_education': 'source', 'what_is_your_current_occupation': 'current_occupation',
                    'what_matters_most_to_you_in_choosing_a_course': 'course_selection',
                    'receive_more_updates_about_our_courses': 'courses_updates',
                    'update_me_on_supply_chain_content': 'supply_chain_content_updates',
                    'get_updates_on_dm_content': 'dm_content_updates',
                    'i_agree_to_pay_the_amount_through_cheque': 'cheque_payment',
                    'a_free_copy_of_mastering_the_interview': 'mastering_interview'}, inplace=True)

df.head(1)
```

| | prospect_id | lead_number | lead_origin | lead_source | do_not_email | do_not_call | converted | total_visits | time_on_website |
|---|--------------------------------------|-------------|-------------|-------------|--------------|-------------|-----------|--------------|-----------------|
| 0 | 7927b2df-8bba-4d29-b9a2-b6e0beafe620 | 660737 | API | Olark Chat | No | No | 0 | 0.00 | |

Drop prospect_id column

```
df.drop('prospect_id', axis = 1, inplace = True)
```

Replace "Select" category with null values

```
# Select all non-numeric columns
df_obj = df.select_dtypes(include='object')

# Find out columns that have "Select"
s = lambda x: x.str.contains('Select', na=False)
l = df_obj.columns[df_obj.apply(s).any()].tolist()
print(l)
```

```
['specialization', 'source', 'lead_profile', 'city']
```

There are 4 columns that contains Select , which are effectively null values. We are going to make that change

```
# select all the columns that have a "Select" entry
sel_cols = ['specialization', 'source', 'lead_profile', 'city']

# replace values
df[sel_cols] = df[sel_cols].replace('Select', np.NaN)
```

Handle null values and sales generated columns

- Given there are a number of columns with very high number of null entries, let's calculate the percentage of null values in each column, and take a decision from there.
- Furthermore, we can also drop Sales generated columns because those are the data entries that are made after the sales team has connected with the student. Those data have no bearing to the purpose of our model ie. providing lead score. The columns are
 - tags
 - lead_quality
 - all asymmetrique columns
 - last_activity
 - last_notable_activity

```
# Calculate percentage of null values for each column
(df.isnull().sum() / df.shape[0]) * 100
```

| | |
|--------------|------|
| lead_number | 0.00 |
| lead_origin | 0.00 |
| lead_source | 0.39 |
| do_not_email | 0.00 |
| do_not_call | 0.00 |
| converted | 0.00 |

```

total_visits          1.48
time_on_website       0.00
page_views_per_visit  1.48
last_activity         1.11
country              26.63
specialization        36.58
source               78.46
occupation           29.11
course_selection_reason 29.32
search               0.00
magazine             0.00
newspaper_article    0.00
x_education_forums   0.00
newspaper            0.00
digital_advertisement 0.00
through_recommendations 0.00
courses_updates      0.00
tags                36.29
lead_quality         51.59
supply_chain_content_updates 0.00
dm_content_updates   0.00
lead_profile        74.19
city                39.71
asymmetrique_activity_index 45.65
asymmetrique_profile_index 45.65
asymmetrique_activity_score 45.65
asymmetrique_profile_score 45.65
cheque_payment       0.00
mastering_interview  0.00
last_notable_activity 0.00
dtype: float64

```

Observation: As can be seen, there are quite a few columns with high number of missing data. Since there are no ways to get data back from reliable sources, we can drop all those columns that have missing values > 40%

Drop columns that have null values > 40% or Sales generated columns

```

df.drop(['source', 'lead_quality', 'lead_profile', 'asymmetrique_activity_index',
        'asymmetrique_profile_index', 'asymmetrique_activity_score', 'asy
        'tags', 'last_activity', 'last_notable_activity'],
        axis = 1, inplace = True)

df.head(1)

```

| | lead_number | lead_origin | lead_source | do_not_email | do_not_call | converted | total_visits | time_on_website | page_view |
|---|-------------|-------------|-------------|--------------|-------------|-----------|--------------|-----------------|-----------|
| 0 | 660737 | API | Olark Chat | No | No | 0 | 0.00 | | 0 |

```

# Lets look at what are we left with
# Calculate percentage of null values for each column

```



```
(df.isnull().sum() / df.shape[0]) * 100
```

| | |
|------------------------------|-------|
| lead_number | 0.00 |
| lead_origin | 0.00 |
| lead_source | 0.39 |
| do_not_email | 0.00 |
| do_not_call | 0.00 |
| converted | 0.00 |
| total_visits | 1.48 |
| time_on_website | 0.00 |
| page_views_per_visit | 1.48 |
| country | 26.63 |
| specialization | 36.58 |
| occupation | 29.11 |
| course_selection_reason | 29.32 |
| search | 0.00 |
| magazine | 0.00 |
| newspaper_article | 0.00 |
| x_education_forums | 0.00 |
| newspaper | 0.00 |
| digital_advertisement | 0.00 |
| through_recommendations | 0.00 |
| courses_updates | 0.00 |
| supply_chain_content_updates | 0.00 |
| dm_content_updates | 0.00 |
| city | 39.71 |
| cheque_payment | 0.00 |
| mastering_interview | 0.00 |
| dtype: float64 | |

Observations

There are five columns that still have high null values: country , specialization , occupation , course_selection_reason , and city . We will look at them individually to see what can be done

country column

```
df.country.value_counts(normalize = True, dropna = False) * 100
```

| | |
|----------------------|-------|
| India | 70.26 |
| NaN | 26.63 |
| United States | 0.75 |
| United Arab Emirates | 0.57 |
| Singapore | 0.26 |
| Saudi Arabia | 0.23 |
| United Kingdom | 0.16 |
| Australia | 0.14 |
| Qatar | 0.11 |
| Bahrain | 0.08 |
| Hong Kong | 0.08 |
| France | 0.06 |

| | |
|---------------------|------|
| Oman | 0.06 |
| unknown | 0.05 |
| South Africa | 0.04 |
| Nigeria | 0.04 |
| Germany | 0.04 |
| Kuwait | 0.04 |
| Canada | 0.04 |
| Sweden | 0.03 |
| Bangladesh | 0.02 |
| Italy | 0.02 |
| Belgium | 0.02 |
| Ghana | 0.02 |
| Uganda | 0.02 |
| China | 0.02 |
| Asia/Pacific Region | 0.02 |
| Philippines | 0.02 |
| Netherlands | 0.02 |
| Switzerland | 0.01 |
| Malaysia | 0.01 |
| Kenya | 0.01 |
| Liberia | 0.01 |
| Sri Lanka | 0.01 |
| Vietnam | 0.01 |
| Tanzania | 0.01 |
| Denmark | 0.01 |
| Indonesia | 0.01 |
| Russia | 0.01 |

Name: country, dtype: float64

Observation

The distribution of the data is very heavily skewed, with India + null values = 97% of the total. It is safe to drop this column.

```
df.drop('country', axis = 1, inplace = True)
```

course_selection_reason column

```
df.course_selection_reason.value_counts(normalize = True, dropna = False) * 100
```

| | |
|---------------------------|-------|
| Better Career Prospects | 70.65 |
| NaN | 29.32 |
| Flexibility & Convenience | 0.02 |
| Other | 0.01 |

Name: course_selection_reason, dtype: float64

Observation

The distribution of the data is very heavily skewed, with Better career prospects + null values = approx 100% of the total. It is safe to drop this column.

```
df.drop('course_selection_reason', axis = 1, inplace = True)
```

occupation column

```
df.occupation.value_counts(normalize = True, dropna = False) * 100
```

```
Unemployed          60.61
NaN                 29.11
Working Professional  7.64
Student              2.27
Other                0.17
Housewife            0.11
Businessman          0.09
Name: occupation, dtype: float64
```

Observation

For occupation, we can first combine categories, and then impute proportionally to maintain the distribution and not introduce bias

```
# combine low representing categories
```

```
df.loc[(df.occupation == 'Student') | (df.occupation == 'Other') | (df.occupation == 'Housewife') |  
       (df.occupation == 'Businessman') , 'occupation'] = 'Student and Others'
```

```
df.occupation.value_counts(normalize = True) * 100
```

```
Unemployed          85.50
Working Professional 10.78
Student and Others    3.73
Name: occupation, dtype: float64
```

```
# impute proportionately
```

```
df['occupation'] = df.occupation.fillna(pd.Series(np.random.choice(['Unemployed', 'Working Professional',  
                                                                    'Student and Others',  
                                                                    'Unemployed'],  
                                                                    size = df.occupation.shape[0],  
                                                                    p = [0.8550, 0.1078, 0.0372])))
```

specialization column

```
df.specialization.value_counts(normalize = True, dropna = False) * 100
```

```
NaN                 36.58
Finance Management  10.56
Human Resource Management  9.18
Marketing Management  9.07
Operations Management  5.44
Business Administration  4.36
IT Projects Management  3.96
Supply Chain Management  3.78
Banking, Investment And Insurance  3.66
```

| | |
|------------------------|------|
| Media and Advertising | 2.20 |
| Travel and Tourism | 2.20 |
| International Business | 1.93 |
| Healthcare Management | 1.72 |
| Hospitality Management | 1.23 |
| E-COMMERCE | 1.21 |
| Retail Management | 1.08 |
| Rural and Agribusiness | 0.79 |
| E-Business | 0.62 |
| Services Excellence | 0.43 |

Name: specialization, dtype: float64

Observation

For specialization, we can first combine categories based on the course type, and then impute proportionally to maintain the distribution and not introduce bias

```
# categorize all management courses
df.loc[(df.specialization == 'Finance Management') | (df.specialization == 'Human Resou
(df.specialization == 'Marketing Management') | (df.specialization == 'Operatic
(df.specialization == 'IT Projects Management') | (df.specialization == 'Supply
(df.specialization == 'Healthcare Management') | (df.specialization == 'Hospital
(df.specialization == 'Retail Management') , 'specialization'] = 'Management Spe

# categorize all busines courses
df.loc[(df.specialization == 'Business Administration') | (df.specialization == 'Intern
(df.specialization == 'Rural and Agribusiness') | (df.specialization == 'E-Busin
, 'specialization'] = 'Business Specializations'

# categorize all industry courses
df.loc[(df.specialization == 'Banking, Investment And Insurance') | (df.specialization
(df.specialization == 'Travel and Tourism') | (df.specialization == 'Services Ex
(df.specialization == 'E-COMMERCE'), 'specialization'] = 'Industry Specializatio
```

```
df.specialization.value_counts(normalize = True) * 100
```

| | |
|----------------------------|-------|
| Management Specializations | 72.58 |
| Industry Specializations | 15.29 |
| Business Specializations | 12.13 |

Name: specialization, dtype: float64

```
# impute proportionately
df['specialization'] = df.specialization.fillna(pd.Series(np.random.choice(['Management
'Business Specializations', 'Indust
p = [0.7258, 0.1213,
```

city column

```
df.city.value_counts(normalize = True, dropna = False) * 100
```

| | |
|-----------------------------|-------|
| NaN | 39.71 |
| Mumbai | 34.87 |
| Thane & Outskirts | 8.14 |
| Other Cities | 7.42 |
| Other Cities of Maharashtra | 4.95 |
| Other Metro Cities | 4.11 |
| Tier II Cities | 0.80 |

Name: city, dtype: float64

Observations We will categorize cities based on logical decisions and impute proportionately

```
# categorize all non-mumbai, but Maharashtra cities
df.loc[(df.city == 'Thane & Outskirts') | (df.city == 'Other Cities of Maharashtra'),
        'city'] = 'Non-Mumbai Maharashtra Cities'

# categorize all other cities
df.loc[(df.city == 'Other Cities') | (df.city == 'Other Metro Cities') | (df.city == 'Tier II Cities'),
        'city'] = 'Non-Maharashtra Cities'
```

```
df.city.value_counts(normalize = True) * 100
```

| | |
|-------------------------------|-------|
| Mumbai | 57.84 |
| Non-Mumbai Maharashtra Cities | 21.70 |
| Non-Maharashtra Cities | 20.46 |

Name: city, dtype: float64

```
# impute proportionately
df['city'] = df.city.fillna(pd.Series(np.random.choice(['Mumbai', 'Non-Mumbai Maharashtra Cities', 'Non-Maharashtra Cities'],
                                                         size = df.city.isnull().sum(),
                                                         p = [0.5784, 0.2170, 0.2046]))
```

Handle categorical columns with low number of missing values and low representation of categories

In this step, we will go through the rest of the categorical columns one by one and

- Merge categories that have low representation
- Impute the missing values

```
(df.isnull().sum() / df.shape[0]) * 100
```

| | |
|--------------|------|
| lead_number | 0.00 |
| lead_origin | 0.00 |
| lead_source | 0.39 |
| do_not_email | 0.00 |
| do_not_call | 0.00 |
| converted | 0.00 |
| total_visits | 1.48 |

| | |
|------------------------------|------|
| time_on_website | 0.00 |
| page_views_per_visit | 1.48 |
| specialization | 0.00 |
| occupation | 0.00 |
| search | 0.00 |
| magazine | 0.00 |
| newspaper_article | 0.00 |
| x_education_forums | 0.00 |
| newspaper | 0.00 |
| digital_advertisement | 0.00 |
| through_recommendations | 0.00 |
| courses_updates | 0.00 |
| supply_chain_content_updates | 0.00 |
| dm_content_updates | 0.00 |
| city | 0.00 |
| cheque_payment | 0.00 |
| mastering_interview | 0.00 |

dtype: float64

```
# determine unique values for all object datatype columns
for k, v in df.select_dtypes(include='object').nunique().to_dict().items():
    print('{} = {}'.format(k,v))
```

```
lead_origin = 5
lead_source = 21
do_not_email = 2
do_not_call = 2
specialization = 3
occupation = 3
search = 2
magazine = 1
newspaper_article = 2
x_education_forums = 2
newspaper = 2
digital_advertisement = 2
through_recommendations = 2
courses_updates = 1
supply_chain_content_updates = 1
dm_content_updates = 1
city = 3
cheque_payment = 1
mastering_interview = 2
```

Observation

As can be seen from the above output, the categorical columns (i.e. number of unique values > 2) are:

- lead_origin

- lead_source

lead_origin column

```
df.lead_origin.value_counts(normalize = True, dropna = False) * 100
```

```
Landing Page Submission    52.88
API                        38.74
Lead Add Form              7.77
Lead Import                0.60
Quick Add Form             0.01
Name: lead_origin, dtype: float64
```

```
#There are a lot of smaller values which will not be used as definitive factors, lets group them
df.loc[(df.lead_origin == 'Lead Import') | (df.lead_origin == 'Quick Add Form') | (df.lead_origin == 'Lead Add Form and Others')
, 'lead_origin'] = 'Lead Add Form and Others'
```

lead_source column

```
df.lead_source.value_counts(normalize = True, dropna = False) * 100
```

```
Google                31.04
Direct Traffic        27.52
Olark Chat            18.99
Organic Search        12.49
Reference              5.78
Welingak Website      1.54
Referral Sites        1.35
Facebook              0.60
NaN                   0.39
bing                  0.06
google                0.05
Click2call            0.04
Press_Release         0.02
Live Chat             0.02
Social Media          0.02
youtubechannel        0.01
Pay per Click Ads     0.01
testone               0.01
NC_EDM                0.01
welearnblog_Home      0.01
blog                  0.01
WeLearn               0.01
Name: lead_source, dtype: float64
```

```
# Lets impute the missing values with the mode of data i.e. clearly 'Google'
df.lead_source.fillna(df.lead_source.mode()[0], inplace=True)
```

```
#There are a lot of smaller values which will not be used as definitive factors, lets g
df['lead_source'] = df['lead_source'].apply(lambda x: x if
                                             ((x== 'Google') | (x=='Direct Traffic') | (
                                              (x=='Organic Search') | (x=='Reference'))
                                             else 'Other Social Sites')
```

Handle Binary columns

- Drop those columns that have significant data imbalance
- Drop all those columns that have only 1 unique entry

```
# determine unique values
for k, v in df.select_dtypes(include='object').nunique().to_dict().items():
    print('{} = {}'.format(k,v))
```

```
lead_origin = 3
lead_source = 6
do_not_email = 2
do_not_call = 2
specialization = 3
occupation = 3
search = 2
magazine = 1
newspaper_article = 2
x_education_forums = 2
newspaper = 2
digital_advertisement = 2
through_recommendations = 2
courses_updates = 1
supply_chain_content_updates = 1
dm_content_updates = 1
city = 3
cheque_payment = 1
mastering_interview = 2
```

Observation

- The following columns can be dropped as they have just 1 unique values
 - magazine
 - course_updates
 - supply_chain_content_updates
 - dm_content_updates
 - cheque_payment

Let's now check the data imbalance for the rest of the columns

```
# select rest of the binary columns in a new dataframe
df_bin = df[['do_not_email', 'do_not_call', 'search', 'newspaper_article', 'x_education_forums', 'newspaper', 'digital_advertisement', 'through_recommendations', 'mastering_python']]

# see value counts for each of the columns
for i in df_bin.columns:
    x = (df_bin[i].value_counts(normalize = True)) * 100
    print(x)
    print()
```

No 92.06

Yes 7.94

Name: do_not_email, dtype: float64

No 99.98

Yes 0.02

Name: do_not_call, dtype: float64

No 99.85

Yes 0.15

Name: search, dtype: float64

No 99.98

Yes 0.02

Name: newspaper_article, dtype: float64

No 99.99

Yes 0.01

Name: x_education_forums, dtype: float64

No 99.99

Yes 0.01

Name: newspaper, dtype: float64

No 99.96

Yes 0.04

Name: digital_advertisement, dtype: float64

No 99.92

Yes 0.08

Name: through_recommendations, dtype: float64

No 68.74

Yes 31.26

Name: mastering_interview, dtype: float64

Observations

Because of heavy data imbalance, we can drop the following columns as well

- do_not_call
- search
- newspaper_article
- x_education_forums
- newspaper
- digital_advertisement
- through_recommendations

```
drop_bin = ['do_not_call', 'search', 'newspaper_article', 'x_education_forums',  
            'newspaper', 'digital_advertisement', 'through_recommendations', 'magazine',  
            'supply_chain_content_updates', 'dm_content_updates', 'cheque_payment']  
  
df.drop(drop_bin, axis = 1, inplace = True)
```

Handle Numerical columns

lead_number column: change datatype

lead_number column is a unique identifier for each leads. Therefore, aggregations won't be of any relevance. We should change it to object

```
df.lead_number = df.lead_number.astype('object')
```

total_visits column

For this column, we need to handle the missing values, and can convert the datatype to integer since visits can't be decimal

```
df.total_visits.fillna(df.total_visits.median(), inplace=True)  
df.total_visits = df.total_visits.astype('int')
```

page_views_per_visit column

Handle missing values

```
df.page_views_per_visit.fillna(df.page_views_per_visit.median(), inplace=True)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9240 entries, 0 to 9239
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   lead_number            9240 non-null   object
1   lead_origin            9240 non-null   object
2   lead_source            9240 non-null   object
3   do_not_email           9240 non-null   object
4   converted              9240 non-null   int64
5   total_visits           9240 non-null   int32
6   time_on_website        9240 non-null   int64
7   page_views_per_visit   9240 non-null   float64
8   specialization         9240 non-null   object
9   occupation             9240 non-null   object
10  city                   9240 non-null   object
11  mastering_interview     9240 non-null   object
dtypes: float64(1), int32(1), int64(2), object(8)
memory usage: 830.3+ KB
```

Exploratory Data Analysis

Numerical columns




```
# Set style
plt.style.use('ggplot')

# See distribution of each of these columns
fig = plt.figure(figsize = (14, 10))
plt.subplot(2, 2, 1)
plt.hist(df.total_visits, bins = 20)
plt.title('Total website visits')

plt.subplot(2, 2, 2)
plt.hist(df.time_on_website, bins = 20)
plt.title('Time spent on website')

plt.subplot(2, 2, 3)
plt.hist(df.page_views_per_visit, bins = 20)
plt.title('Average number of page views per visit')

plt.show()
```




Observations

- High peaks and skewed data. There might be a possibility of outliers. We will check them next

Heatmap

```
plt.figure(figsize = (14,12))
sns.heatmap(df[['total_visits', 'time_on_website', 'page_views_per_visit']].corr(), cmap=
plt.show()
```



Observations: No significant correlation such that columns can be dropped

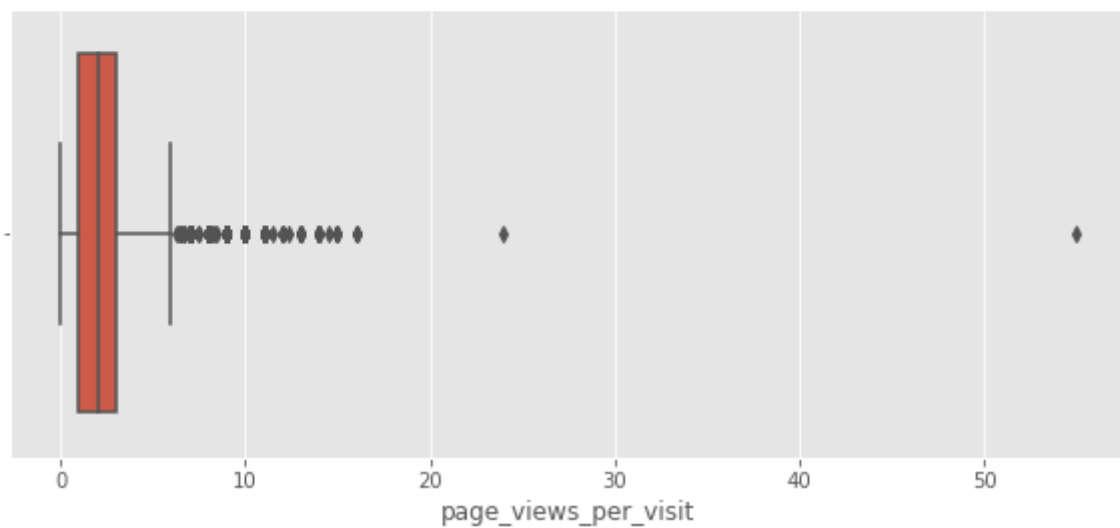


Check for outliers

```
plt.figure(figsize = (10, 14))

plt.subplot(3,1,1)
sns.boxplot(df.total_visits)

plt.subplot(3,1,2)
sns.boxplot(df.time_on_website)

plt.subplot(3,1,3)
sns.boxplot(df.page_views_per_visit)
plt.show()
```




Observations

- Looking at both the box plots and the statistics, there are upper bound outliers in both `total_visits` and `page_views_per_visit` columns. We can also see that the data can be capped at 99 percentile.

Categorical columns

Lead Origin


[illegible]

```
df.head(1)
```

| | lead_number | lead_origin | lead_source | do_not_email | converted | total_visits | time_on_website | page_views_per_visit |
|---|-------------|-------------|-------------|--------------|-----------|--------------|-----------------|----------------------|
| 0 | 660737 | API | Olark Chat | No | 0 | 0.00 | 0 | 0.00 |

Lead Source

```
plt.figure(figsize = (14, 8))  
  
df.groupby('lead_source')['lead_number'].count().sort_values(ascending = False).plot(kind='bar',  
                                                                 edgecolor = 'black',  
                                                                 color = plt.cm.Paired(np.arange(10)),  
                                                                 legend = True,  
                                                                 title = "Lead Number by Lead Source")  
  
plt.show()
```




Specialization

```
plt.figure(figsize = (10, 8))

df.groupby('specialization')['lead_number'].count().sort_values(ascending = False).plot(
    edgecolor = 'black',
    color = plt.cm.Paired(np.arange(3)),
    style = 'b',
    legend = False)


plt.show()
```



Most of the specialization taken are management

Occupation


```
plt.figure(figsize = (14, 8))  
  
df.groupby('occupation')['lead_number'].count().sort_values(ascending = False).plot(kind='bar',  
                                                                 edgecolor = 'black',  
                                                                 color = plt.cm.Paired(np.arange(len(df['occupation'])).astype(int)))  
  
plt.show()
```



Unemployed users are the most significant leads

City

```
plt.figure(figsize = (14, 8))  
  
df.groupby('city')['lead_number'].count().sort_values(ascending = False).plot(kind= 'bar'  
                                             edgecolor = 'black',  
                                             color = plt.cm.Paired(np.arange(  
plt.show()
```




Mumbai in particular and Maharashtra in general dominates the lead. This is likely due to the fact that the courses are based in Mumbai

```
plt.figure(figsize = (14, 8))

df.groupby('do_not_email')['lead_number'].count().sort_values(ascending = False).plot(kind='bar',
    edgecolor = 'black',
    color = plt.cm.Paired(np.arange(2)))

plt.show()
```



Data Preparation

Converting Binary (Yes/No) to 0/1

```
# determine unique values
for k, v in df.select_dtypes(include='object').nunique().to_dict().items():
    print('{} = {}'.format(k,v))
```

```
lead_number = 9240
lead_origin = 3
lead_source = 6
do_not_email = 2
specialization = 3
occupation = 3
city = 3
mastering_interview = 2
```

We have two binary columns: `do_not_email`, `mastering_interview`

```
binlist = ['do_not_email', 'mastering_interview']

# Defining the map function
def binary_map(x):
    return x.map({'Yes': 1, "No": 0})

# Applying the function to the housing list
df[binlist] = df[binlist].apply(binary_map)
```

```
# check the operation was success
df.head()
```

| | lead_number | lead_origin | lead_source | do_not_email | converted | total_visits | time_on_website | page_views_per_visit |
|---|-------------|-------------------------|----------------|--------------|-----------|--------------|-----------------|----------------------|
| 0 | 660737 | API | Olark Chat | 0 | 0 | 0.00 | 0 | 0.00 |
| 1 | 660728 | API | Organic Search | 0 | 0 | 5.00 | 674 | 2.50 |
| 2 | 660727 | Landing Page Submission | Direct Traffic | 0 | 1 | 2.00 | 1532 | 2.00 |
| 3 | 660719 | Landing Page Submission | Direct Traffic | 0 | 0 | 1.00 | 305 | 1.00 |
| 4 | 660681 | Landing Page Submission | Google | 0 | 1 | 2.00 | 1428 | 1.00 |

Creating dummy variable for categorical columns

Categorical columns are: lead_origin , lead_source , specialization , occupation , city

```
# Creating a dummy variable for some of the categorical variables and dropping the first
dummy1 = pd.get_dummies(df[['lead_origin', 'lead_source', 'specialization', 'occupation'])

# Adding the results to the master dataframe
df = pd.concat([df, dummy1], axis=1)
```

```
# Dropping the columns for which dummies have been created
df.drop(['lead_origin', 'lead_source', 'specialization', 'occupation', 'city'], axis = 1)

df.head()
```

Outliers Treatment

```
num_cols = df[['total_visits', 'time_on_website', 'page_views_per_visit']]

# Checking outliers at 25%, 50%, 75%, 90%, 95% and 99%
num_cols.describe(percentiles=[.25, .5, .75, .90, .95, .99])
```

| | total_visits | time_on_website | page_views_per_visit |
|-------|--------------|-----------------|----------------------|
| count | 9240.00 | 9240.00 | 9240.00 |
| mean | 3.44 | 487.70 | 2.36 |
| std | 4.82 | 548.02 | 2.15 |
| min | 0.00 | 0.00 | 0.00 |
| 25% | 1.00 | 12.00 | 1.00 |
| 50% | 3.00 | 248.00 | 2.00 |

| | total_visits | time_on_website | page_views_per_visit |
|-----|--------------|-----------------|----------------------|
| 75% | 5.00 | 936.00 | 3.00 |
| 90% | 7.00 | 1380.00 | 5.00 |
| 95% | 10.00 | 1562.00 | 6.00 |
| 99% | 17.00 | 1840.61 | 9.00 |
| max | 251.00 | 2272.00 | 55.00 |



```
# capping at 99 percentile
```

```
df.total_visits.loc[df.total_visits >= df.total_visits.quantile(0.99)] = df.total_visits.quantile(0.99)
df.page_views_per_visit.loc[df.page_views_per_visit >= df.page_views_per_visit.quantile(0.99)] = df.page_views_per_visit.quantile(0.99)
```

```
plt.figure(figsize = (10, 14))
```

```
plt.subplot(2,1,1)
sns.boxplot(df.total_visits)
```

```
plt.subplot(2,1,2)
sns.boxplot(df.page_views_per_visit)
plt.show()
```



As we can see, we were able to significantly reduce the number of outliers by capping

Test-Train Split

```
# Putting feature variable to X
X = df.drop(['lead_number', 'converted'], axis=1)

X.head(1)
```

| do_not_email | total_visits | time_on_website | page_views_per_visit | mastering_interview | lead_origin_Landing Page Submission | lead_origin Add For |
|--------------|--------------|-----------------|----------------------|---------------------|--|------------------------|
| 0 | 0 | 0.00 | 0 | 0.00 | 0 | 0 |

```
# Putting response variable to y
y = df['converted']

y.head(1)
```

```
0    0
Name: converted, dtype: int64
```

```
# Splitting the data into train and test
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.7, test_size=0.3)
```

Feature Scaling

```
scaler = StandardScaler()

X_train[['total_visits', 'time_on_website', 'page_views_per_visit']] = scaler.fit_transform(
    X_train[['total_visits', 'time_on_website', 'page_views_per_visit']])

X_train.head()
```

| | do_not_email | total_visits | time_on_website | page_views_per_visit | mastering_interview | lead_origin_Landing Page Submission | lead_o Add |
|------|--------------|--------------|-----------------|----------------------|---------------------|--|---------------|
| 1871 | 0 | -1.02 | -0.89 | -1.18 | 0 | 0 | |
| 6795 | 0 | 0.21 | 0.01 | -0.50 | 1 | 1 | |
| 3516 | 0 | 0.51 | -0.69 | 0.09 | 0 | 0 | |
| 8105 | 0 | 0.51 | 1.37 | 1.36 | 0 | 1 | |
| 3934 | 0 | -1.02 | -0.89 | -1.18 | 0 | 0 | |

```
# checking the conversion rate
conversion = (sum(df['converted'])/len(df['converted'].index))*100
conversion
```

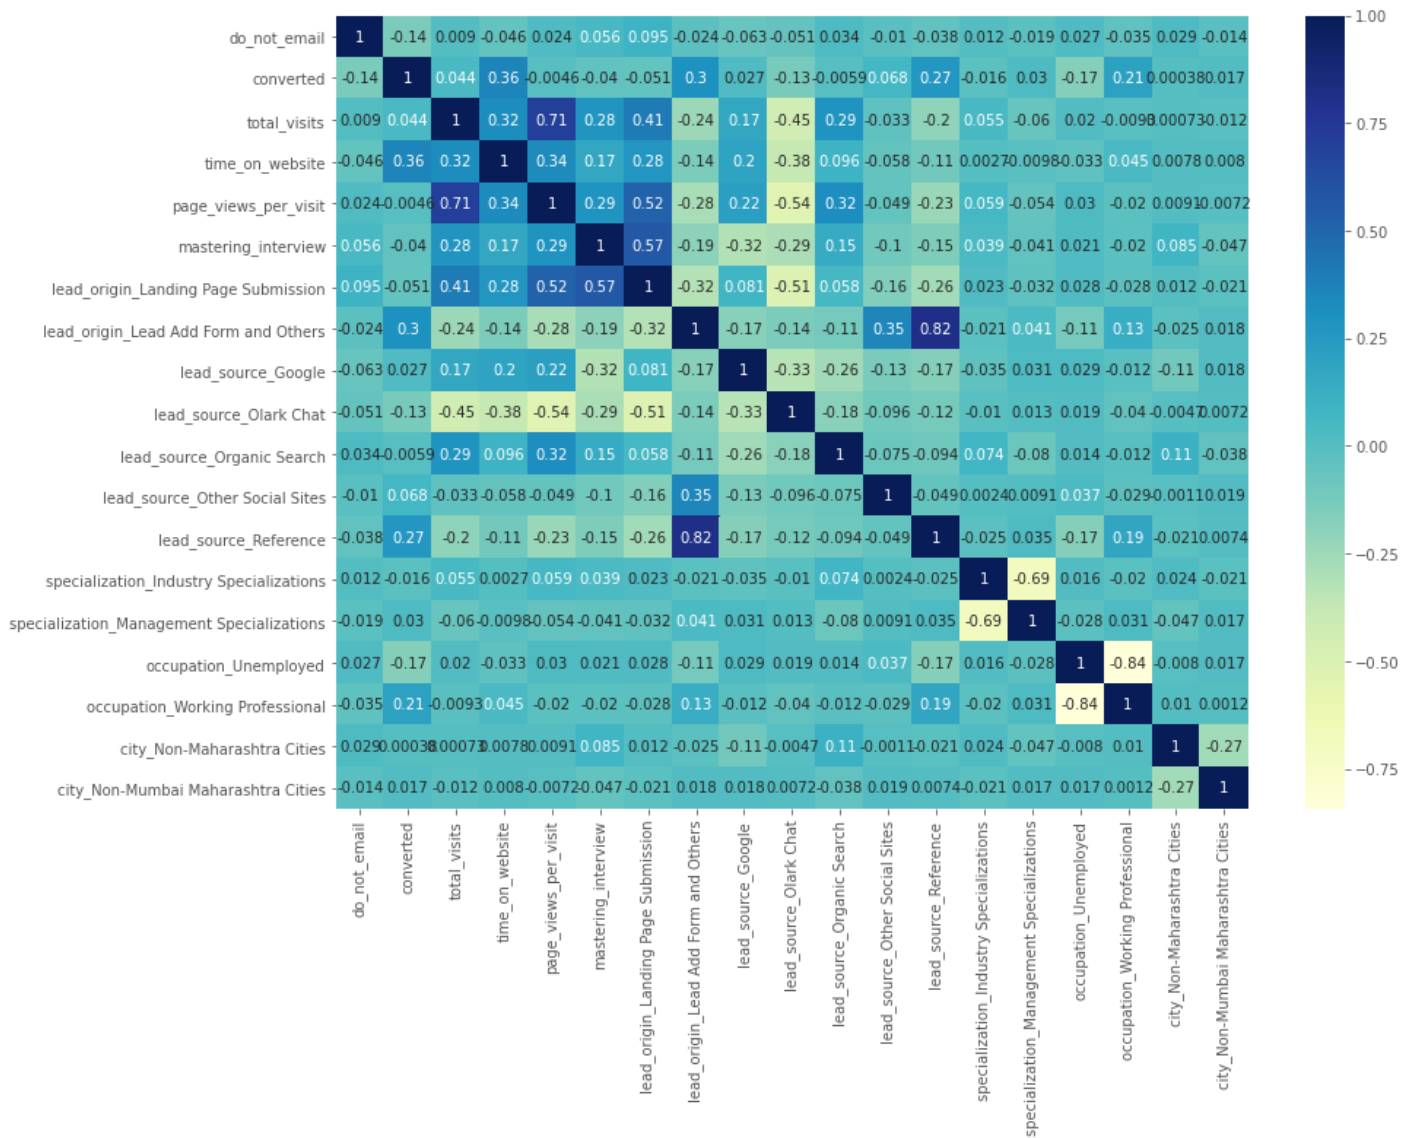
```
38.53896103896104
```

The conversion rate is 38.5%

Looking at correlations

```
# Let's see the correlation matrix
plt.figure(figsize = (14,10))
```

```
sns.heatmap(df.corr(),annot = True, cmap="YlGnBu")
plt.show()
```

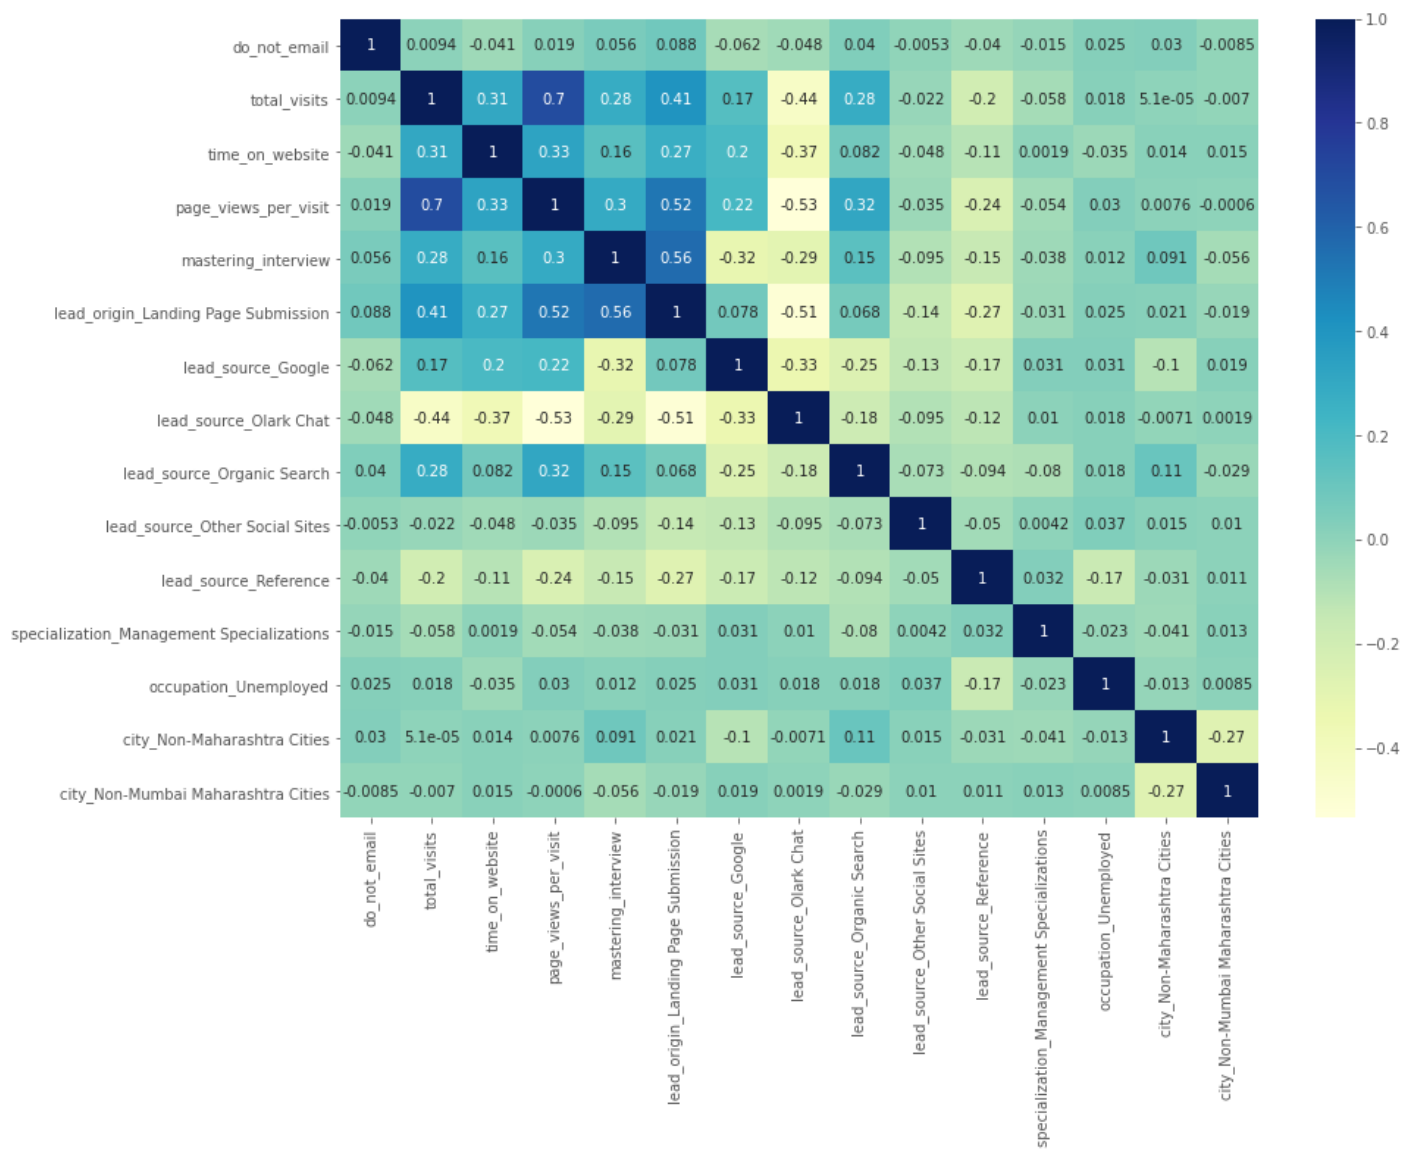


Drop highly correlated dummy variables

```
X_test.drop(['lead_origin_Lead Add Form and Others', 'specialization_Industry Specializations', 'occupation_Working Professional'], axis = 1, inplace = True)
```

```
X_train.drop(['lead_origin_Lead Add Form and Others', 'specialization_Industry Specializations', 'occupation_Working Professional'], axis = 1, inplace = True)
```

```
## lets check the correlation matrix again
plt.figure(figsize = (14,10))
sns.heatmap(X_train.corr(),annot = True, cmap="YlGnBu")
plt.show()
```

Model Building

Model 1: All variables

```
# Logistic regression model
logm1 = sm.GLM(y_train,(sm.add_constant(X_train)), family = sm.families.Binomial())
logm1.fit().summary()
```

Generalized Linear Model Regression Results

| | | | |
|------------------|------------------|-------------------|----------|
| Dep. Variable: | converted | No. Observations: | 6468 |
| Model: | GLM | Df Residuals: | 6452 |
| Model Family: | Binomial | Df Model: | 15 |
| Link Function: | logit | Scale: | 1.0000 |
| Method: | IRLS | Log-Likelihood: | -3304.7 |
| Date: | Sun, 06 Dec 2020 | Deviance: | 6609.4 |
| Time: | 05:01:08 | Pearson chi2: | 6.64e+03 |
| No. Iterations: | 5 | | |
| Covariance Type: | nonrobust | | |

| | coef | std err | z | P> z | [0.025 | 0.975] |
|---|---------|---------|--------|-------|--------|--------|
| const | -0.4513 | 0.147 | -3.080 | 0.002 | -0.739 | -0.164 |
| do_not_email | -1.2167 | 0.144 | -8.467 | 0.000 | -1.498 | -0.935 |
| total_visits | 0.1403 | 0.042 | 3.356 | 0.001 | 0.058 | 0.222 |
| time_on_website | 1.0456 | 0.036 | 29.388 | 0.000 | 0.976 | 1.115 |
| page_views_per_visit | -0.1787 | 0.048 | -3.696 | 0.000 | -0.273 | -0.084 |
| mastering_interview | -0.0088 | 0.094 | -0.094 | 0.925 | -0.192 | 0.175 |
| lead_origin_Landing Page Submission | -0.0012 | 0.092 | -0.013 | 0.990 | -0.182 | 0.180 |
| lead_source_Google | 0.3639 | 0.100 | 3.625 | 0.000 | 0.167 | 0.561 |
| lead_source_Olark Chat | 0.6743 | 0.137 | 4.935 | 0.000 | 0.406 | 0.942 |
| lead_source_Organic Search | 0.2207 | 0.116 | 1.903 | 0.057 | -0.007 | 0.448 |
| lead_source_Other Social Sites | 1.6136 | 0.175 | 9.202 | 0.000 | 1.270 | 1.957 |
| lead_source_Reference | 3.9674 | 0.221 | 17.945 | 0.000 | 3.534 | 4.401 |
| specialization_Management Specializations | 0.1367 | 0.069 | 1.974 | 0.048 | 0.001 | 0.272 |
| occupation_Unemployed | -0.8398 | 0.085 | -9.889 | 0.000 | -1.006 | -0.673 |
| city_Non-Maharashtra Cities | 0.1526 | 0.078 | 1.948 | 0.051 | -0.001 | 0.306 |
| city_Non-Mumbai Maharashtra Cities | 0.0745 | 0.077 | 0.972 | 0.331 | -0.076 | 0.225 |

Feature selection using RFE

```
# initiate logistic regression
logreg = LogisticRegression()

# initiate rfe
rfe = RFE(logreg, 13)          # running RFE with 13 variables as output
rfe = rfe.fit(X_train, y_train)
```

```
rfe.support_
```

```
array([ True,  True,  True,  True, False, False,  True,  True,  True,
        True,  True,  True,  True,  True,  True])
```

```
list(zip(X_train.columns, rfe.support_, rfe.ranking_))
```

```
[('do_not_email', True, 1),
 ('total_visits', True, 1),
 ('time_on_website', True, 1),
 ('page_views_per_visit', True, 1),
 ('mastering_interview', False, 2),
 ('lead_origin_Landing Page Submission', False, 3),
 ('lead_source_Google', True, 1),
 ('lead_source_Olark Chat', True, 1),
 ('lead_source_Organic Search', True, 1),
 ('lead_source_Other Social Sites', True, 1),
 ('lead_source_Reference', True, 1),
```

```
( 'specialization_Management Specializations', True, 1),
( 'occupation_Unemployed', True, 1),
( 'city_Non-Maharashtra Cities', True, 1),
( 'city_Non-Mumbai Maharashtra Cities', True, 1)]
```

```
# assign columns
col = X_train.columns[rfe.support_]
```

```
# check what columns were not selected by RFE
X_train.columns[~rfe.support_]
```

```
Index(['mastering_interview', 'lead_origin_Landing Page Submission'], dtype='object')
```

Model 2: Assessing the model with statsmodel

```
X_train_sm = sm.add_constant(X_train[col])
logm2 = sm.GLM(y_train,X_train_sm, family = sm.families.Binomial())
res = logm2.fit()
res.summary()
```

| Generalized Linear Model Regression Results | | | | | | | |
|---|------------------|-------------------|----------|-------|--------|--------|--|
| Dep. Variable: | converted | No. Observations: | 6468 | | | | |
| Model: | GLM | Df Residuals: | 6454 | | | | |
| Model Family: | Binomial | Df Model: | 13 | | | | |
| Link Function: | logit | Scale: | 1.0000 | | | | |
| Method: | IRLS | Log-Likelihood: | -3304.7 | | | | |
| Date: | Sun, 06 Dec 2020 | Deviance: | 6609.4 | | | | |
| Time: | 05:08:43 | Pearson chi2: | 6.64e+03 | | | | |
| No. Iterations: | 5 | | | | | | |
| Covariance Type: | nonrobust | | | | | | |
| | coef | std err | z | P> z | [0.025 | 0.975] | |
| const | -0.4592 | 0.112 | -4.082 | 0.000 | -0.680 | -0.239 | |
| do_not_email | -1.2163 | 0.144 | -8.474 | 0.000 | -1.498 | -0.935 | |
| total_visits | 0.1399 | 0.042 | 3.364 | 0.001 | 0.058 | 0.221 | |
| time_on_website | 1.0456 | 0.036 | 29.393 | 0.000 | 0.976 | 1.115 | |
| page_views_per_visit | -0.1791 | 0.047 | -3.811 | 0.000 | -0.271 | -0.087 | |
| lead_source_Google | 0.3702 | 0.079 | 4.675 | 0.000 | 0.215 | 0.525 | |
| lead_source_Olark Chat | 0.6807 | 0.113 | 6.032 | 0.000 | 0.460 | 0.902 | |
| lead_source_Organic Search | 0.2240 | 0.108 | 2.068 | 0.039 | 0.012 | 0.436 | |
| lead_source_Other Social Sites | 1.6203 | 0.158 | 10.283 | 0.000 | 1.311 | 1.929 | |
| lead_source_Reference | 3.9740 | 0.207 | 19.210 | 0.000 | 3.569 | 4.379 | |
| specialization_Management Specializations | 0.1368 | 0.069 | 1.975 | 0.048 | 0.001 | 0.273 | |
| occupation_Unemployed | -0.8397 | 0.085 | -9.889 | 0.000 | -1.006 | -0.673 | |

| | | | | | | |
|------------------------------------|--------|-------|-------|-------|--------|-------|
| city_Non-Maharashtra Cities | 0.1524 | 0.078 | 1.946 | 0.052 | -0.001 | 0.306 |
| city_Non-Mumbai Maharashtra Cities | 0.0749 | 0.077 | 0.979 | 0.328 | -0.075 | 0.225 |

Lets create empty lists of categorical columns and numerical columns, then we will review the columns one by one and see what needs to be done with each of them

```
cat= []
num = []
```

```
#Creating a function to get the column details
def details(x):
    print(df[x].value_counts())
    print(df[x].isnull().sum(), 'null values')
    print(df[x].isnull().sum()/df.count()['lead_number']*100, '% values are null')
```

```
# Do not email column
details('Do Not Email')
```

we'll leave this column as is for now, and add this to a new list of binary categorical variables

```
bi_cat = []
bi_cat.append('Do Not Email')
```

```
#Do Not Call column
details('Do Not Call')
```

```
#Lets drop the column since its only two records and it doesn't make sense to keep this
df.drop('Do Not Call',axis=1,inplace = True)
#df[['Do Not Call', 'Converted']].value_counts()
```

```
details('TotalVisits')
```

```
# We can see that there is a lot of outliers here, we can also plot a boxplot to get a
sns.boxplot(df['TotalVisits'])
#and see the median(sometimes we impute the null values with median), and 95th to 99th
df.TotalVisits.quantile([0.50,0.95,0.96,0.97,0.98,0.99])
```

Considering the values and outliers here, Let's cap the values at 96th percentile i.e. 10

Also, we will impute the null values with 0 here, we could impute this with median but considering the values might not have been tracked because they haven't logged on to the site, and 0 is also the mode of the column.

```
df[df['TotalVisits'].isnull()][ 'TotalVisits' ] = df['TotalVisits'].mode()[0]
#df[df.TotalVisits > df.TotalVisits.quantile(0.96)] = df.TotalVisits.quantile(0.96)
#Also add this column to our numerical columns list
num.append('TotalVisits')
```

Let's also create a function to cap the outliers if needed in future

```
def cap(col,typ='right',value=0.95):
    if typ == 'left':
        df[df[col]<df[col].quantile(value)][col] = df[col].quantile(value)
    else:
        df[df[col]>df[col].quantile(value)][col] = df[col].quantile(value)
```

```
# and capping the column as mentioned earlier
cap('TotalVisits')
```

```
# Lets look at Total Time Spent on Website column
details('Total Time Spent on Website')
```

```
#####
##Also add this column to our numerical columns list
#num.append('Total Time Spent on Website')
```

Since we also look at the boxplots and percentiles for numerical numbers, lets create a similar details function to include these two pieces of information.

```
def num_details(x):
    print(df[x].value_counts())
    print(df[x].isnull().sum(), 'null values')
    print(df[x].isnull().sum()/df.count()[x]*100, '% values are null')
    print('Percentiles are as follows')
    print(df[x].quantile([0.50,0.95,0.96,0.97,0.98,0.99]))
    sns.boxplot(df[x])
```

```
#Lets look at column Page Views Per Visit
num_details('Page Views Per Visit')
```

```
df[df['Page Views Per Visit'].isnull()]['Page Views Per Visit'] = df['Page Views Per Vi  
cap('Page Views Per Visit')  
#Also add this column to our numerical columns list  
num.append('Page Views Per Visit')
```

```
details('Last Activity')
```

```
#####  
##Also add this column to our numerical columns list  
#num.append('Last Activity')
```

```
#Lets have a look at the column 'Country'  
details('Country')
```

There are 27 percent null values in this column, and other values are mostly India, so this column would not be that useful, we could create a binary column using this like 'India' if there weren't so many null values here. but considering the null values, lets drop this column

```
df.drop('Country', axis =1, inplace = True)
```

```
details('Specialization')
```

```
#####
```

```
details('How did you hear about X Education')
```

There are too many null values in this as well, and the most values are also not selected I think, because the option says default value 'Select', lets drop this column as well

```
df.drop('How did you hear about X Education', axis =1, inplace = True)
```

```
details('What is your current occupation')
```

```
#####
```

```
details('What matters most to you in choosing a course')
```

```
#####
```

```
details('Search')
```

```
details('Magazine')
```

```
details('Newspaper Article')
```

```
details('X Education Forums')
```

```
details('Newspaper')
```

```
details('Digital Advertisement')
```

All these columns have high imbalance, and mostly have only one value i.e. No. So it doesn't make sense to keep these columns, Let's delete these columns

```
#Let's look at the column 'Through Recommendation'  
details('Through Recommendations')
```

```
details('Receive More Updates About Our Courses')
```

```
df.drop(['Search', 'Magazine', 'Newspaper Article', 'X Education Forums', 'Newspaper', 'Digi
```

```
details('Tags')
```

```
#####
```

```
details('Lead Quality')
```

```
#####
```

```
details('Update me on Supply Chain Content')
```

```
# There's only singular value in this, so let's drop this column  
df.drop('Update me on Supply Chain Content', axis=1, inplace = True)
```

```
details('Get updates on DM Content')
```

```
# There's only singular value in this, so let's drop this column  
df.drop('Get updates on DM Content', axis=1, inplace = True)
```

```
details('Lead Profile')
```

```
#####
```

```
details('City')
```

```
#####
```

```
details('Asymmetrique Activity Index')
```

```
#####
```

```
details('Asymmetrique Profile Index')
```

```
#####
```

```
details('Asymmetrique Activity Score')
```

```
#####
```

```
details('Asymmetrique Profile Score')
```

```
#####
```

```
details('I agree to pay the amount through cheque')
```

```
# There's only singular value in this, so lets drop this column  
df.drop('I agree to pay the amount through cheque', axis=1, inplace = True)
```

```
details('A free copy of Mastering The Interview')
```

we'll leave this column as is for now, and add this to a new list of binary categorical variables

```
bi_cat.append('A free copy of Mastering The Interview')
```

```
details('Last Notable Activity')
```

```
#####
```

```
#Lets have a look at our three categories of column  
print(cat)
```



```
print(num)  
print(bi_cat)
```

```
df.head(1)
```